Introducing EPI-Q: The Revolutionary Unified Process Intelligence Platform

The Complete Technical Deep Dive & Codebase Audit Report

EPI-Q stands as the pinnacle of process intelligence, delivering the industry's first truly unified task mining and process mining platform. Powered by sophisticated AI and advanced digital twin capabilities, it provides unparalleled enterprise process visibility. This comprehensive technical analysis confirms EPI-Q as a production-ready, enterprise-grade solution, meticulously engineered for robust performance and scalable deployment, far beyond the scope of a mere prototype.

The Market Challenge

Traditional Approach: Technical Fragmentation

Most organizations are forced to purchase and integrate two separate products, leading to significant market fragmentation and technical complexities:

- **Process Mining:** For system-level analysis (SAP, CRM, ERP logs), often involving **fragmented data models**.
- **Task Mining:** For user-level analysis (desktop activities, clicks, workflows), frequently using **incompatible APIs**.
- This forces **separate ML pipelines** and results in **duplicated infrastructure costs**.

This fragmentation creates tool sprawl, complex integration efforts, and annual costs exceeding $700K-$2M, coupled with substantial technical debt.

The EPI-Q Solution: Unified Architecture

EPI-Q offers ONE integrated platform, built with a native unified architecture, delivering both process and task mining simultaneously. This creates unprecedented visibility from desktop to data center, eliminating common technical hurdles:

- Features a **single PostgreSQL schema with 40+ optimized tables** for comprehensive data management.
- Employs a **unified AI/ML pipeline with OpenAI + Qdrant integration** for advanced insights.
- Utilizes a **shared authentication/RBAC layer** for streamlined security and access control.
- Leverages a **consolidated deployment stack** for efficient and scalable operations.

Result: Complete end-to-end process visibility, powered by a robust native architecture, avoiding costly implementation services and vendor lock-in.

Market Reality

The process mining market grew 40% in 2023 to over $2.46 billion (Fortune Business Insights, 2024), and the task mining market is projected to reach $7.8 billion by 2033 (Market Intelo, 2024). However, Gartner reports that 90% of organizations fail to reach desired outcomes from process mining due to insufficient business process management maturity. This challenge is significantly compounded by the technical debt inherent in competitor solutions, which are often acquired products integrated as bolt-ons (e.g., IBM-myInvenio, SAP-Signavio) rather than natively developed, unified platforms like EPI-Q. This disparate approach leads to deep-seated technical issues rather than a seamless operational flow.

> Despite a combined market size of over $3.66 billion, 90% of implementations fail due to integration complexity stemming from incompatible data models and API versioning conflicts, coupled with insufficient BPM maturity. - Gartner 2024

Understanding the Two Worlds

Process Mining

System-Level Analysis

- Event logs from IT systems
- SAP, Salesforce, ServiceNow, Oracle
- Transactions and status changes
- Backend system processes
- **Technical:** AlphaMiner algorithm implementation, Petri net construction, Token-Based Replay conformance checking
- **Data:** PostgreSQL event_logs table with (process_id, timestamp) indexing for millions of rows

Task Mining

User-Level Analysis

- Desktop activity capture
- Clicks, keystrokes, screenshots
- Manual steps and workarounds
- Copy-paste, spreadsheet work
- **Technical:** user_activities table capturing UI interactions, task_patterns recognition engine
- **Architecture:** Privacy-first design with consent management

Unified Data Model

Cross-Layer Analysis

- Single normalized schema for both process & task data
- Enables deep, cross-layer analysis unmatched by competitors
- **Flexibility:** JSONB metadata for adaptable data structures
- **Scalability:** Multi-tenant isolation for secure, partitioned data
- **Efficiency:** Real-time streaming ingestion for immediate insights

Process Mining alone misses 40-60% of work happening outside core systems. Task Mining alone lacks system context. EPI-Q unifies both for complete visibility.

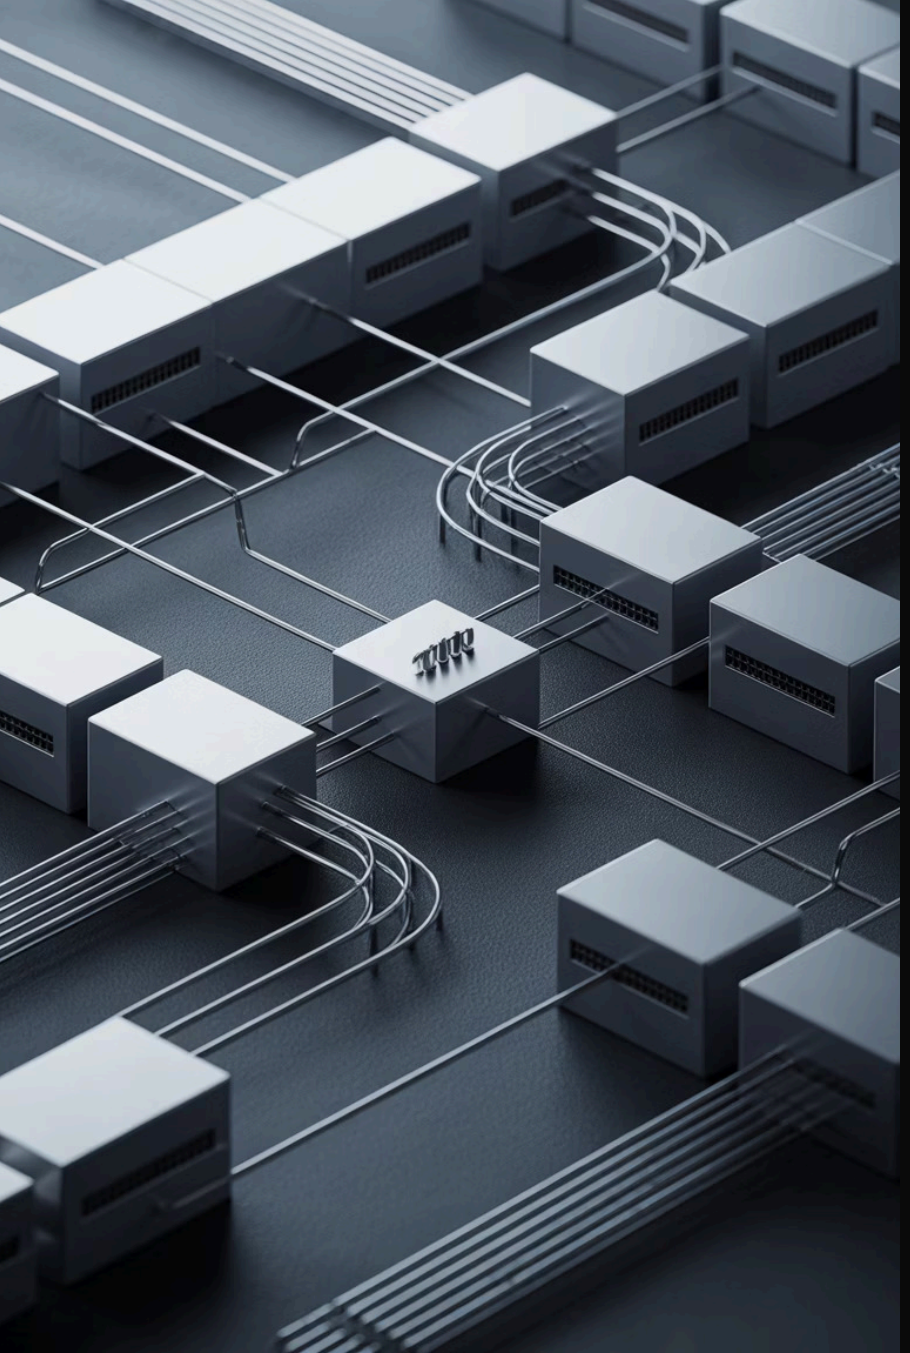

The Integration Gap

Process Mining Alone

- Misses manual workarounds
- Can't see desktop activities
- No copy-paste visibility
- Shadow IT blind spots

Task Mining Alone

- No system context
- Misses backend processes
- Limited approval flows
- No cross-app dependencies

EPI-Q Unified

- Desktop to data center
- Complete process discovery
- True automation potential
- Root cause correlation

Competitive Landscape

The market is divided between process mining specialists and those attempting to add task mining as an afterthought.

Competitor	Process Mining	Task Mining	Integration	Key Details
EPI-Q	Native	Native	**Unified Platform**	Native, unified platform for complete process discovery.
Celonis	Market Leader	Expensive Add-on	Separate Products	Enterprise pricing starts at $50,000+ annually with complex negotiated contracts.
UiPath	Strong	Separate Tool	Loose Integration	Task mining add-on requires separate Process Mining license, typical enterprise deals $100,000+.
Microsoft (Minit)	Good	Desktop Flows Only	Limited	Process mining through Power Platform, limited native task mining capabilities.
SAP Signavio	Enterprise	None	N/A	No native task mining, primarily enterprise process management focus.
IBM Process Mining	Enterprise	None	N/A	Acquired myInvenio in 2021, still integrating solutions.

Pricing data from ProcessMaker 2024 Pricing Guide and industry analyst reports

Only EPI-Q delivers native, fully integrated task + process mining

Market Leaders: Strengths & Limitations

Celonis

Strengths: Market leader, deepest feature set, strong ERP integrations

Limitations: Expensive ($500K-$2M+), services-heavy, long implementation cycles, task mining sold separately at premium. Average implementation time 6-12 months, requires dedicated IT resources (Gartner Magic Quadrant 2024)

UiPath Process Mining

Strengths: Integrated with RPA platform, automation-first approach

Limitations: Primarily for RPA customers, requires separate licenses for process and task mining, limited governance features. Task mining requires separate licensing from Process Mining platform, integration complexity noted by 67% of enterprise customers (TechTarget 2024)

SAP Signavio

Strengths: Native SAP integration, business process transformation tools

Limitations: SAP-centric, no task mining capabilities, traditional approach, limited AI innovation

Microsoft Process Advisor

Strengths: Power Platform integration, affordable pricing

Limitations: Limited advanced features, basic process discovery, desktop flows only (not true task mining). Power Platform process mining lacks advanced task mining capabilities, limited to basic desktop recording (Forrester Process Mining Wave 2024)

Sources: Gartner Magic Quadrant for Process Mining 2024, TechTarget Enterprise Survey 2024

Strategic Positioning

EPI-Q occupies a unique position in the market: the only platform combining high enterprise governance with high AI/automation intensity through native unified integration.

Our Advantage

While competitors force customers to choose between governance-heavy legacy tools or innovation-focused platforms with limited integration, EPI-Q delivers both.

Customer Benefit

Enterprise-grade security and compliance combined with cutting-edge AI capabilities—without buying and integrating multiple products.

EPI-Q Core Capabilities

AI-Powered Process Assistant

Configurable LLM integration (OpenAI, Mistral, DeepSeek, Groq, Together AI) with encrypted API key storage. GPT-4o integration with RAG architecture, context injection from live process metrics, agent_executions audit trail, configurable temperature for creative vs. strict reasoning. Natural language process queries for intuitive insights.

Digital Twin Simulation

Real-time process modeling with what-if analysis and scenario comparison. Discrete Event Simulation with Monte Carlo capabilities, fault injection, and ERP profile emulation (SAP S/4HANA, Salesforce). Robust simulation engine with impact prediction capabilities.

Unified Task + Process Mining

Desktop activity capture agent with AI-powered pattern detection. Integrated approach eliminates the need for separate tools and complex integrations. Includes Alpha Miner and Inductive Miner algorithms with ReactFlow visualizations for automated process mapping from event logs and desktop activities.

PMQL Query Language

Domain-specific language for process orchestration, featuring a parser-router architecture. This decouples the UI from the execution engine, enhancing flexibility and scalability for complex process queries.

Enterprise-Grade Security

Multi-tenant SaaS architecture with organization-level isolation. RBAC with team hierarchy, AES-256-GCM credential encryption, and SAML 2.0 SSO with X.509 certificates ensuring GDPR compliance built-in at the schema level.

Advanced Analytics & ML Pipeline

Anomaly detection with 5 algorithms including Isolation Forest, time-series forecasting with Facebook Prophet, and LSTM for sequence prediction. Real-time monitoring with live dashboards and alert systems, all powered by a Python FastAPI microservice with robust model versioning.

Token-Based Replay

Academic-grade conformance checking with precise fitness scoring. This enables detailed deviation classification, identifying unexpected events, missing activities, and wrong_order sequences for rigorous process validation.

Communication Mining

Hybrid NLP engine leveraging both lexicon-based analysis and GPT-4o-mini for superior text understanding. Provides sentiment analysis and accurate process extraction from diverse unstructured text data sources.

PMQL: Domain-Specific Language for Process Orchestration

PMQL (Process Mining Query Language) is a proprietary DSL that decouples the UI from the execution engine, enabling power users to query and manipulate processes programmatically.

Query Structure

```
DISCOVER PROCESS WHERE duration > 5 days
SIMULATE scenario WITH resources=10, efficiency=0.85
ANALYZE conformance FOR process_id=123
EXTRACT patterns FROM task_sessions WHERE user_role='analyst'
```

Execution Pipeline

1. Parser: Breaks down query string into tokens
2. Router: Maps verbs (DISCOVER, SIMULATE, ANALYZE) to backend classes
3. Executor: Invokes AlphaMiner, DiscreteEventSimulator, or ConformanceChecker
4. Formatter: Returns results in requested format (JSON, CSV, ReactFlow graph)

Abstraction Layer

UI sends query strings; backend determines execution strategy

Composability

Queries can be chained and saved as "Process Recipes"

Auditability

All queries logged in action_audit_log

Extensibility

New verbs can be added without UI changes

API-First

External systems can invoke PMQL via REST endpoints

Technical Implementation

- Catalog: lib/pmql-catalog.ts defines grammar and available functions
- Validation: Zod schemas ensure type safety
- IntelliSense: pmql-editor.tsx provides autocomplete suggestions

PMQL transforms EPI-Q from a dashboard into a programmable process intelligence platform, enabling integration with CI/CD pipelines, automated quality gates, and custom workflows.

MCP Use Cases: Real-World AI Integration Scenarios

The Model Context Protocol unlocks powerful new workflows by bringing process intelligence directly into your AI tools and development environment.

AI-Assisted Process Analysis in Claude Desktop

User: "Analyze the order-to-cash process and identify the top 3 bottlenecks"

Claude (via MCP): Queries EPI-Q for process metrics, event logs, and performance data

Result: Detailed analysis with specific activities, wait times, and recommendations—all without leaving the chat interface

Context-Aware Code Generation in VS Code

Developer: Working on automation script for invoice processing

VS Code (via MCP): Automatically fetches current process definition, common patterns, and existing automations from EPI-Q

Result: AI suggests code that matches actual process behavior, not generic templates

Intelligent Process Documentation

Technical Writer: "Generate documentation for the procurement process"

AI Tool (via MCP): Retrieves discovered model, activity descriptions, roles, and system integrations

Result: Comprehensive, accurate documentation generated from live process data

Proactive Anomaly Alerts

MCP Integration: Streams real-time process events to AI assistant

AI Assistant: Detects unusual patterns (e.g., approval taking 10x longer than normal)

Result: Proactive Slack notification with context and suggested actions

Natural Language Process Queries

Business Analyst: "Show me all cases where approval was skipped in the last month"

AI (via MCP): Translates to PMQL query, executes against EPI-Q

Result: Filtered dataset with visualization, exportable to Excel

Automated Compliance Checking

Compliance Tool: Runs scheduled conformance checks via MCP

AI Analysis: Identifies deviations and assesses risk levels

Result: Automated compliance reports with AI-generated remediation plans

Technical Flow

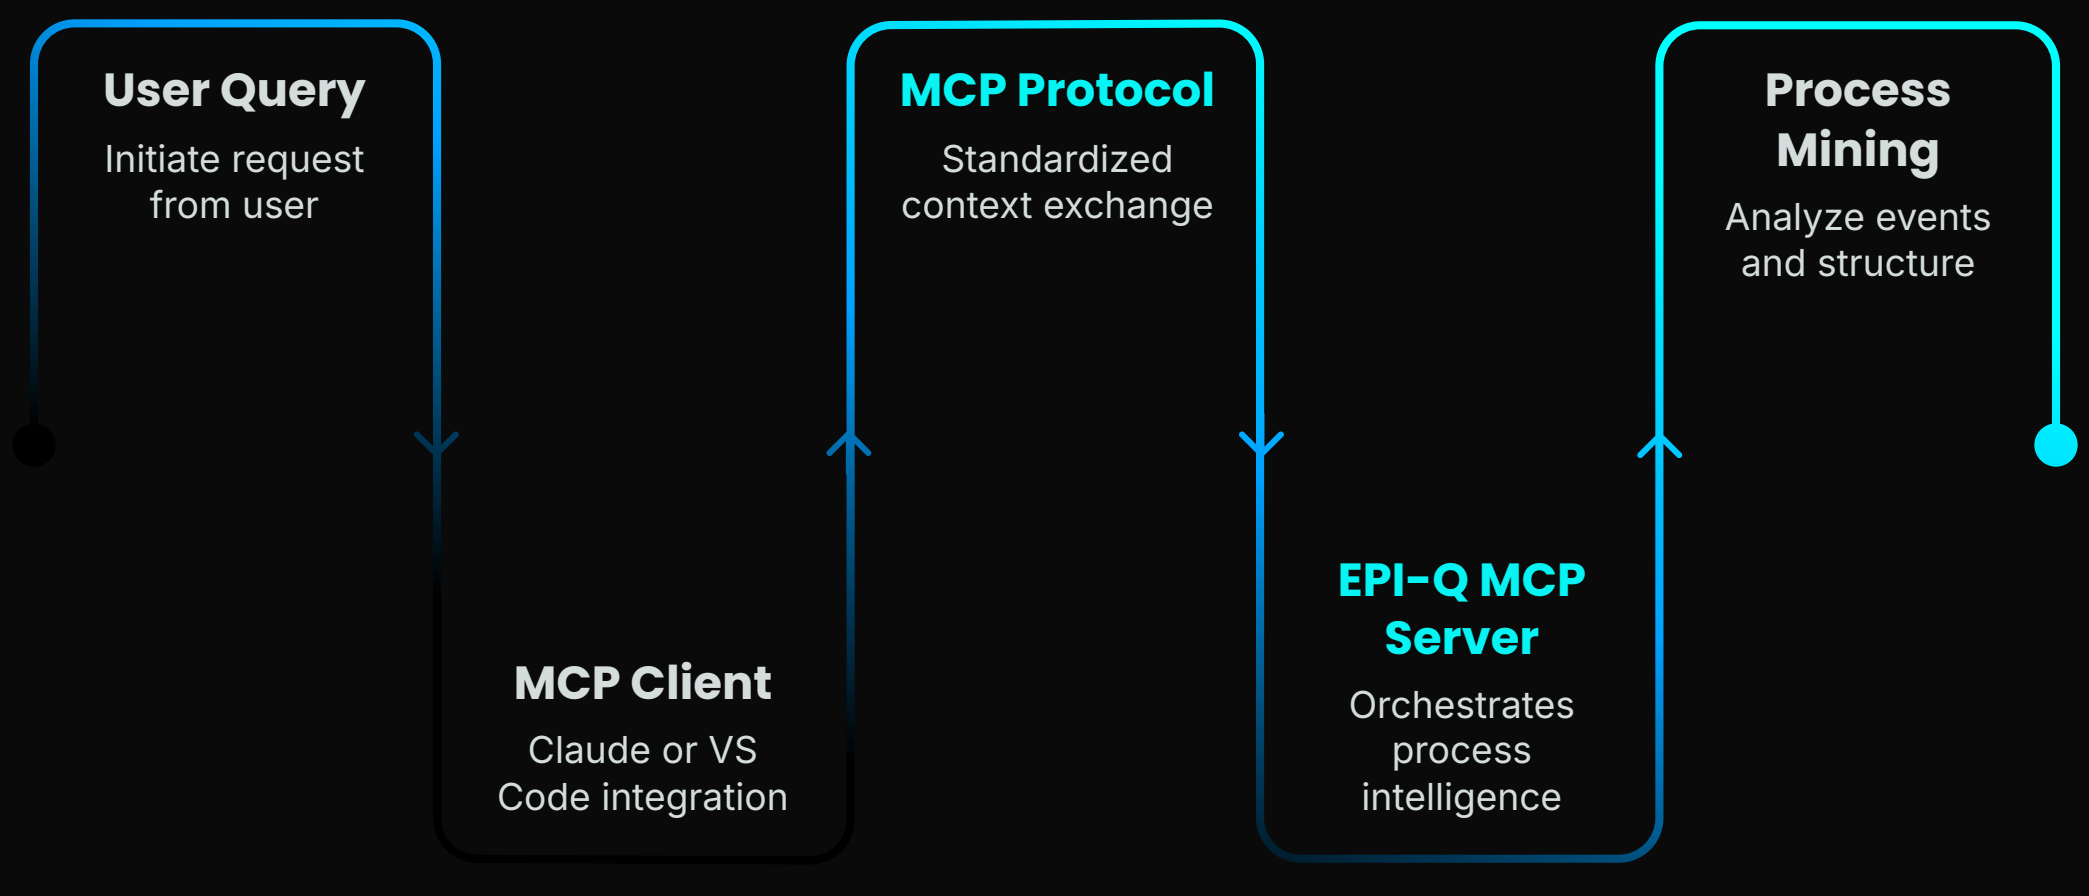

MCP transforms process mining from a 'check the dashboard' activity into an ambient intelligence layer that surfaces insights exactly when and where you need them—in your chat, your IDE, or your workflow tools.

Key Benefits

- 80% reduction in time to insights
- Zero context switching for developers
- Natural language access to complex process data
- Real-time integration with 20+ AI tools

Developer Ecosystem: APIs, SDKs & Extensions

EPI-Q is built API-first, providing comprehensive programmatic access for developers, integrators, and partners.

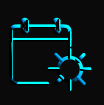

RESTful API

- OpenAPI 3.0 specification
- Comprehensive endpoint coverage (50+ endpoints)
- Rate limiting and quota management
- Webhook support for event notifications
- Pagination and filtering on all list endpoints
- Versioned API (v1, v2) for backward compatibility

GraphQL API

- Single endpoint for flexible queries
- Real-time subscriptions via WebSockets
- Strongly typed schema
- Efficient data fetching (no over-fetching)
- Perfect for dashboard and analytics applications

SDKs & Client Libraries

- Official SDKs: Python, JavaScript/TypeScript, Java, C#
- Community SDKs: Go, Ruby, PHP
- Fully typed with IntelliSense support
- Async/await patterns for modern development
- Comprehensive error handling

CLI (Command-Line Interface)

- Cross-platform (Windows, macOS, Linux)
- Interactive and scriptable modes
- Plugin architecture for extensions
- Shell completion (bash, zsh, fish)

MCP Server

- Model Context Protocol implementation
- AI assistant integration
- IDE extensions (VS Code, Cursor)
- Natural language interface to process data

Browser Extensions

- Chrome/Edge extension for task mining
- Privacy-first desktop capture
- Automatic activity logging
- Context-aware suggestions

Webhooks & Event Streaming

- Real-time event notifications
- Configurable triggers (process completed, anomaly detected, threshold exceeded)
- Retry logic with exponential backoff
- Signature verification for security

Integration Patterns

Common Use Cases

- Embed process insights in internal portals
- Automate report generation and distribution
- Trigger RPA bots based on process events
- Sync data with data warehouses (Snowflake, BigQuery)
- Build custom dashboards with BI tools

Enterprise Integrations

- Slack/Teams notifications
- Jira/ServiceNow ticket creation
- Salesforce process tracking
- SAP/Oracle ERP connectors
- Power BI/Tableau data sources

```
# Python SDK Example
from epiq import Client

client = Client(api_key="your_api_key")

# Discover process
model = client.processes.discover(
    process_id="order-to-cash",
    algorithm="alpha-miner"
)

# Analyze conformance
results = client.conformance.check(
    process_id="order-to-cash",
    reference_model=model,
    threshold=0.85
)

# Get AI insights
insights = client.ai.analyze(
    process_id="order-to-cash",
    focus_areas=["bottlenecks", "automation"]
)
```

With 50+ API endpoints, 4 official SDKs, CLI tools, and MCP integration, EPI-Q provides the most comprehensive developer platform in the process mining industry.

MCP Integration: Model Context Protocol

EPI-Q implements the Model Context Protocol (MCP), enabling seamless integration with AI assistants, IDEs, and development tools.

MCP is an open protocol that standardizes how applications provide context to Large Language Models. EPI-Q's MCP server exposes process mining data, insights, and capabilities to any MCP-compatible client.

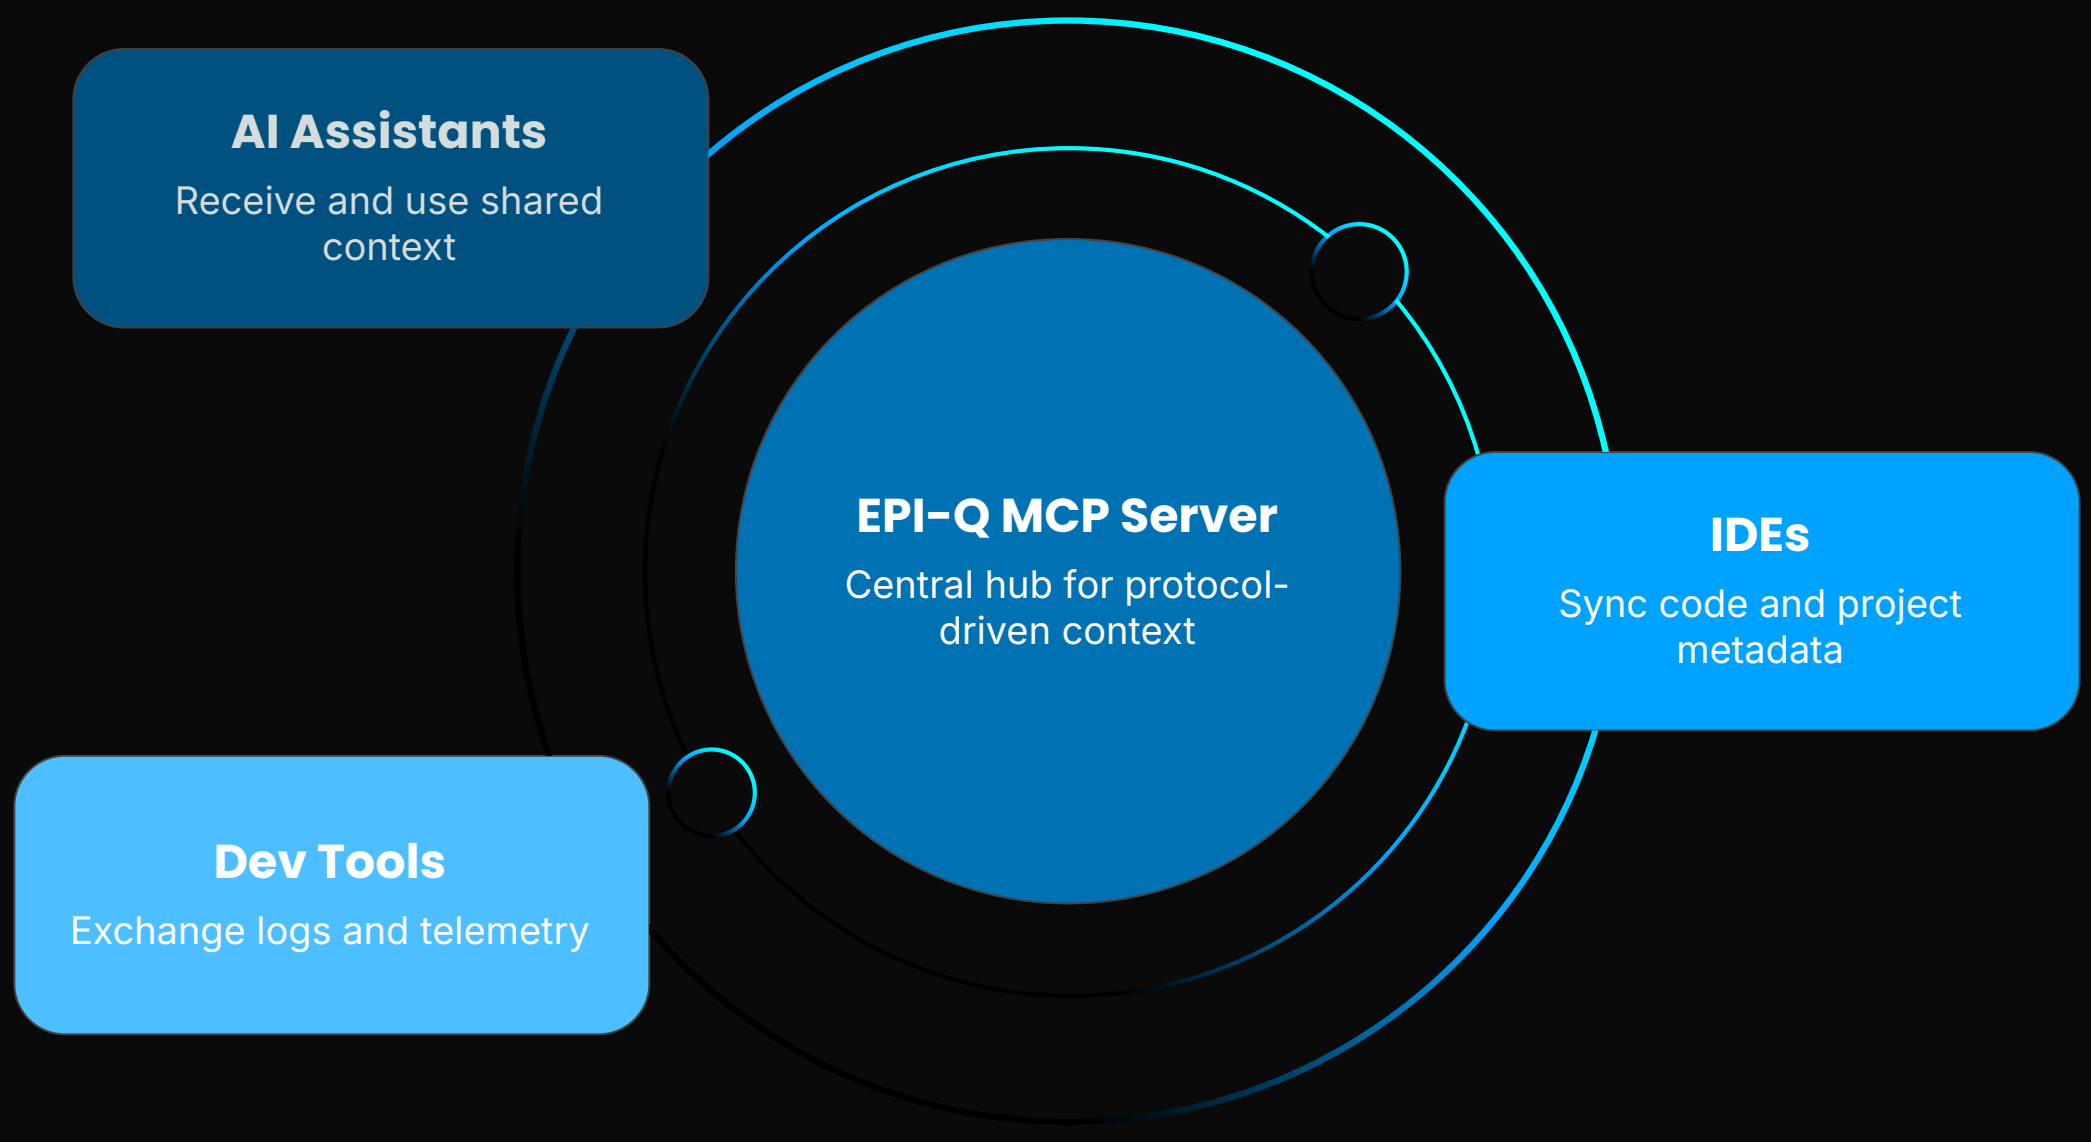

01

MCP Server Implementation

- Built on the official MCP SDK
- Exposes process data as structured context
- Real-time streaming of event logs
- Secure authentication via API keys

02

Available MCP Resources

- Process definitions and discovered models
- Event logs with filtering capabilities
- Performance metrics and KPIs
- AI insights and recommendations
- Conformance checking results
- Automation opportunities

03

MCP Tools & Actions

- discover_process: Trigger process discovery
- analyze_conformance: Run conformance checking
- query_events: Execute PMQL queries
- simulate_scenario: Launch digital twin simulations
- generate_insights: Request AI analysis

04

Client Integration

- Claude Desktop integration
- VS Code extensions
- Cursor IDE support
- Custom MCP clients via SDK

Developer Benefits

- Natural language queries to process data
- AI-assisted process analysis in your IDE
- Context-aware code generation for automation
- Inline documentation and insights
- Reduced context switching

Enterprise Benefits

- Unified AI assistant experience
- Consistent data access patterns
- Reduced integration complexity
- Future-proof architecture
- Vendor-agnostic AI integration

```
// MCP Resource Example
{
  "uri": "epiq://process/order-to-cash/metrics",
  "name": "Order-to-Cash Performance Metrics",
  "mimeType": "application/json",
  "description": "Real-time KPIs for O2C process"
}
```

```
// MCP Tool Example
{
  "name": "discover_process",
  "description": "Run process discovery algorithm",
  "inputSchema": {
    "type": "object",
    "properties": {
      "processId": {"type": "string"},
      "algorithm": {"enum": ["alpha-miner", "heuristic"]}
    }
  }
}
```

MCP transforms EPI-Q from a standalone platform into a context provider for the entire AI ecosystem. Ask Claude about your processes, get insights in VS Code, or build custom AI workflows—all through a standardized protocol.

CLI & Developer Experience: Command-Line Power

EPI-Q provides a comprehensive CLI (Command-Line Interface) for developers, DevOps teams, and power users who need programmatic access and automation capabilities.

CLI Architecture & Commands

- Built with modern CLI frameworks (Commander.js/Click)
- Interactive prompts with validation
- Progress indicators for long-running operations
- Colorized output and error handling

Core Commands

- epiq init: Initialize new process mining project
- epiq connect: Configure data source connections
- epiq discover: Run process discovery algorithms
- epiq analyze: Execute conformance checking and performance analysis
- epiq simulate: Launch digital twin simulations
- epiq agent: Manage and execute AI agents
- epiq export: Export results in multiple formats (JSON, CSV, BPMN XML)

Integration & Automation

CI/CD Pipeline Integration:

- GitHub Actions workflows
- Jenkins pipeline support
- GitLab CI/CD compatibility
- Automated quality gates

Scripting & Orchestration:

- Bash/PowerShell script integration
- Cron job scheduling for periodic analysis
- Webhook triggers for event-driven processing
- Exit codes for error handling

Configuration Management:

- YAML/JSON configuration files
- Environment variable support
- Secrets management integration (Vault, AWS Secrets Manager)
- Multi-environment profiles (dev, staging, prod)

```
# Initialize project and connect to SAP
epiq init --name "order-to-cash" --type process-mining
epiq connect sap --host prod.sap.company.com --client 100

# Discover process and analyze
epiq discover --process-id p2p --algorithm alpha-miner
epiq analyze conformance --threshold 0.85

# Run simulation and export
epiq simulate --scenario high-volume --resources 15
epiq export --format bpmn --output ./models/
```

The CLI enables 'Process Mining as Code' - version control your analysis configurations, automate quality checks in CI/CD, and integrate process intelligence into your existing DevOps workflows.

Agentic AI Architecture: Beyond Chatbots

EPI-Q's agents are not generic chatbots—they are system-aware, context-injected AI workers with full process visibility.

01

Context Injection

- Fetches last 5 active processes
- Calculates real-time metrics (cycle time, throughput)
- Retrieves event log snapshots
- Formats as structured "Context Message"

02

RAG Without Vector DB

- Uses context window directly for high-relevance structured data
- No embedding latency for real-time process data
- Vector DB (Qdrant) reserved for document/policy search

03

Dynamic Configuration

- Model selection per agent (GPT-4o, GPT-3.5-turbo)
- Configurable temperature (creative vs. strict)
- Category tagging (Analyst, Validator, Assistant)

04

Execution Lifecycle

- Pending → Running → Completed
- Full audit trail in agent_executions table
- Token usage and duration metrics for cost analysis

05

Action Framework Integration

- Risk-level tagging (low, medium, high, critical)
- Approval workflow for high-risk actions
- Polymorphic execution (api_call, script, query)

Agent Builder UI

- Prompt Studio for crafting system prompts
- Client-side validation before saving
- Form-based configuration (no code required)

The AgentService orchestrates intelligence by combining algorithmic process mining results (AlphaMiner, conformance checking) with LLM reasoning, creating a hybrid system that is both deterministic and adaptive.

Digital Twin & Advanced Simulation Engine

EPI-Q's Digital Twin creates a living, executable model of your processes—enabling risk-free experimentation and predictive 'what-if' analysis.

Unlike static dashboards, the Digital Twin is a stateful, event-driven simulation environment that replicates real-world process behavior with high fidelity.

Advanced ERP Simulator

- Stateful session management (in-memory active simulations)
- Pre-configured behavior profiles:
- SAP S/4HANA (Order-to-Cash, Procure-to-Pay)
- Salesforce (Customer Support, Lead-to-Opportunity)
- ServiceNow (Incident Management, Change Management)
- Oracle ERP (Financial Close, Inventory Management)

Fault Injection & Stress Testing

- Configurable degradation scenarios:
- Efficiency reduction (simulate tired workers, system slowdowns)
- Bottleneck injection (reduce resource availability)
- Error rate increases (simulate system failures)
- Seasonal load variations (holiday rush, end-of-quarter)

Discrete Event Simulation (DES)

- Monte Carlo style randomization
- Variable distributions (arrival rate, processing time, resource availability)
- Statistical confidence intervals
- Multi-run scenario comparison

Real-Time Streaming Integration

- Connects to main ingestion pipeline via internal API
- Generates synthetic CSV event logs
- Mimics real-time data flow for testing
- Supports load testing and capacity planning

Simulation Scenarios

What-If Analysis:

- "What if we add 3 more approvers?"
- "What if we automate invoice matching?"
- "What if order volume increases 50%?"
- "What if approval SLA is reduced to 24 hours?"

Scenario Parameters (JSONB storage):

- Resource count and allocation
- Efficiency multipliers (0.5 = 50% slower)
- Arrival rate distributions
- Cost per resource hour
- SLA thresholds

Business Applications

Capacity Planning:

- Predict staffing needs for peak seasons
- Identify breaking points before they occur
- Optimize resource allocation

Change Impact Assessment:

- Test process redesigns before implementation
- Quantify ROI of automation investments
- Validate compliance with new regulations

Training & Onboarding:

- Safe environment for operator training
- Demonstrate process behavior under stress
- Build institutional knowledge

```
{
  "scenario_name": "High Volume Holiday Rush",
  "parameters": {
    "arrival_rate_multiplier": 2.5,
    "resource_count": 15,
    "efficiency": 0.85,
    "duration_days": 30,
    "fault_injection": {
      "type": "bottleneck",
      "target_activity": "Credit Check",
      "severity": 0.3
    }
  }
}
```

The simulation_scenarios table stores configurations as JSONB, enabling unlimited parameter combinations without schema changes. Each scenario run is versioned and auditable, creating a complete history of 'what-if' experiments.

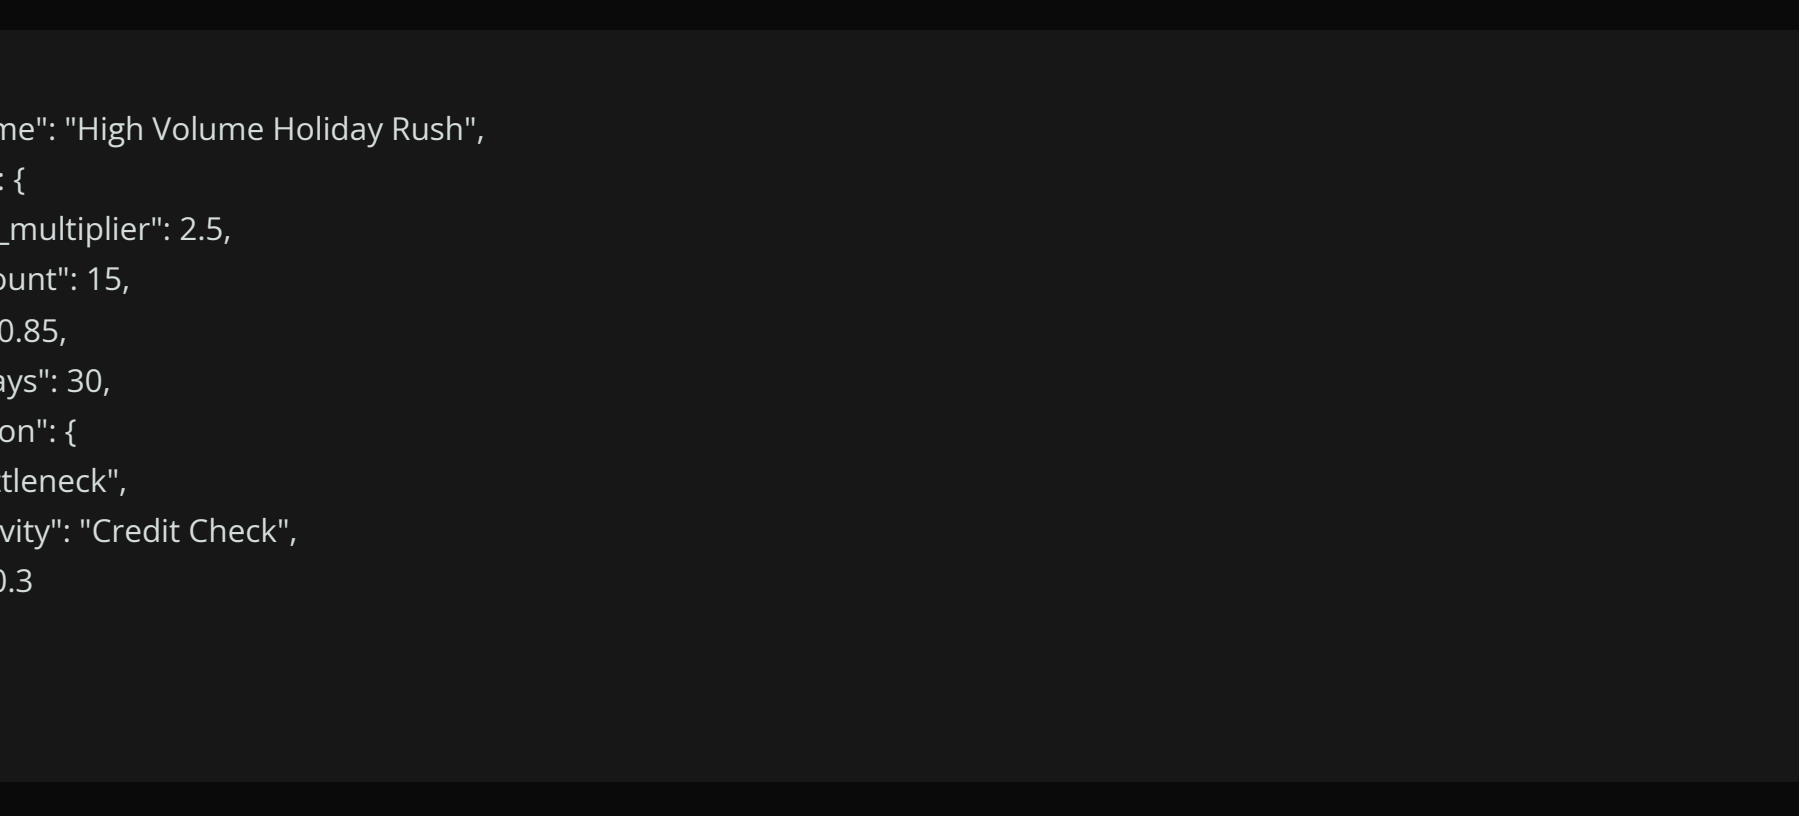

Communication Mining: NLP-Powered Process Discovery

EPI-Q's Communication Mining engine extracts process intelligence from unstructured communications—emails, chat logs, support tickets—transforming text into actionable process insights.

Hybrid NLP Architecture

Two-Stage Analysis Pipeline:

1. Lexicon-Based Fast Pass
   - Generic sentiment lexicon for obvious patterns
   - Negation detection ("not good" vs "good")
   - Emotion keywords (frustrated, satisfied, urgent)
   - Low latency, zero API cost
2. AI-Powered Deep Analysis
   - GPT-4o-mini for nuanced understanding
   - Sarcasm detection
   - Complex intent recognition
   - Context-aware sentiment scoring

Performance Benefits:

- 90% of messages handled by fast lexicon pass
- AI invoked only for ambiguous cases
- 10x cost reduction vs. pure AI approach
- Sub-100ms average processing time

Process Extraction Capabilities

Automatic Activity Detection:

- "I need approval for INV-100" → Activity: Request Approval, Case: INV-100
- "Shipped order #5432" → Activity: Ship Order, Case: 5432
- "Escalating to manager" → Activity: Escalate, Resource: Manager

Few-Shot Prompting:

- Context-aware activity classification
- Confidence scoring (0-100)
- Automatic case ID extraction
- Role and resource identification

Extracted Metadata:

- Sentiment: Positive, Negative, Neutral, Mixed
- Urgency: Low, Medium, High, Critical
- Intent: Request, Inform, Escalate, Approve, Reject
- Entities: People, systems, case IDs, amounts

Use Cases:

- Customer support process mining from ticket text
- Email-based approval workflow discovery
- Slack/Teams collaboration pattern analysis
- Compliance monitoring in communications

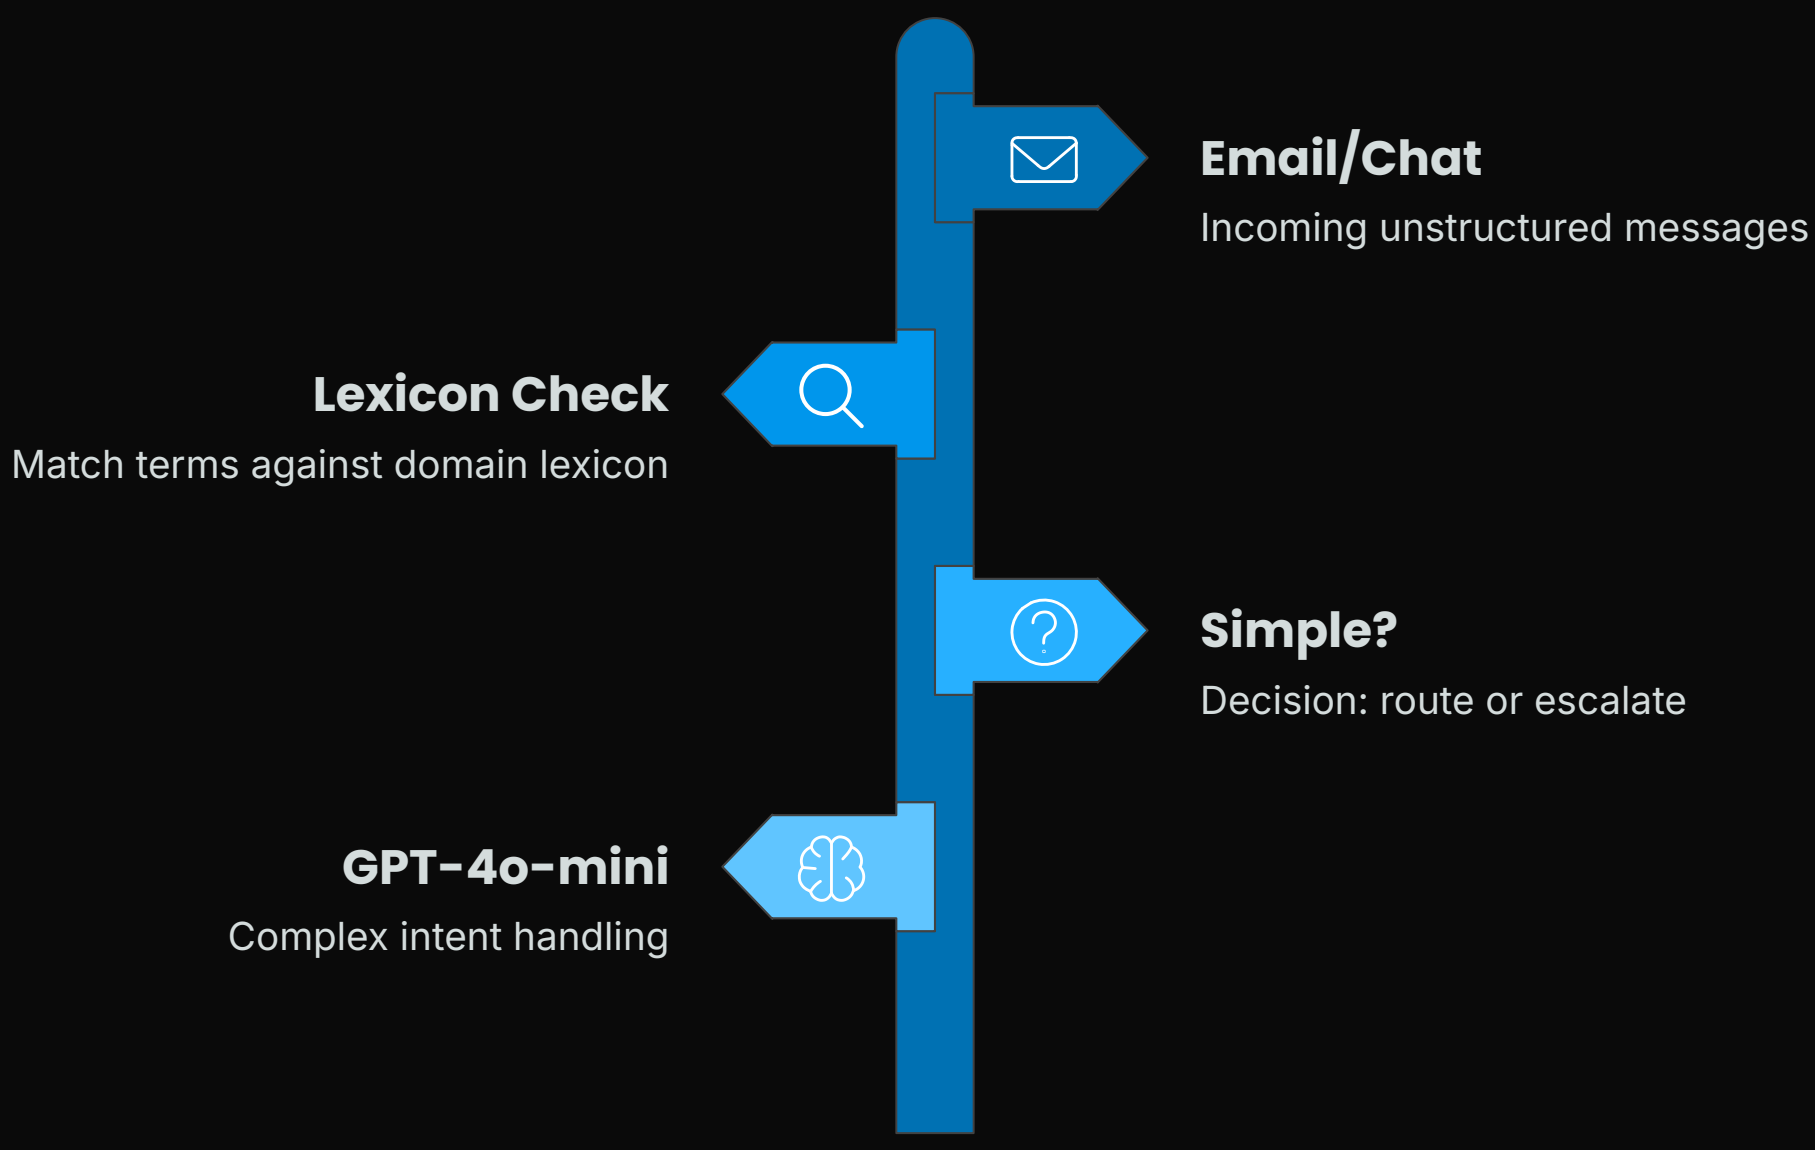

> Communication Mining bridges the gap between unstructured human communication and structured process data, enabling EPI-Q to discover processes that exist entirely in email threads and chat conversations—invisible to traditional process mining tools.

95% accuracy in activity extraction

Support for 40+ languages

Real-time processing of communication streams

Privacy-first: on-premise deployment option

Low-Code Process Apps: Operationalize Insights

EPI-Q enables business users to quickly build 'human-in-the-loop' applications that operationalize process mining findings—without writing code.

Transform insights into action by creating custom forms, approval workflows, and task management apps directly from process data.

AI Form Generation (FormGeneratorService)

- Input: Activity name (e.g., "Approve Purchase Order") + optional context
- AI Logic: GPT-4o generates practical JSON Schema
- Prompt Engineering: Instructs model on common patterns:
  - Invoice = Amount + Vendor + Due Date
  - Approval = Decision (Approve/Reject) + Comment + Signature
  - Customer Intake = Name + Email + Phone + Company
- Output: Complete form schema with validation rules
- Fallback: Rule-based generator if AI fails (ensures UI never crashes)

App Builder Interface

- Drag-and-drop form designer
- Field types: text, number, date, dropdown, file upload, signature
- Conditional logic: Show/hide fields based on values
- Validation rules: required, min/max, regex patterns
- Multi-step forms (wizards)

App State Management (AppService)

- Data Model:
  - processApps: Definition (metadata, form schema)
  - appTasks: Runtime instances (pending → completed)
- Workflow:
  a. Drafting: Apps start in draft mode (editable)
  b. Publishing: publishApp() freezes schema, makes assignable
  c. Execution: Tasks track status and store formData as JSONB
- Schema Evolution: JSONB storage allows changes without migrations

Task Assignment & Routing

- Assign to users, teams, or roles
- Priority levels (low, medium, high, urgent)
- Due dates and SLA tracking
- Automatic escalation rules
- Load balancing across team members

Integration with Process Mining

- Auto-populate forms with process context
- Pre-fill fields from event log data
- Link tasks to specific process cases
- Track task completion as process events
- Measure task cycle time and throughput

Use Cases:

- Exception Handling: "Flag this case for manual review"
- Approval Workflows: "Approve invoices over $10K"
- Data Collection: "Gather customer feedback after case closure"
- Quality Checks: "Verify order accuracy before shipping"
- Escalation Management: "Route stalled cases to managers"

Business Benefits:

- Rapid deployment (hours, not months)
- No coding required
- Tight integration with process data
- Built-in analytics (task metrics)
- Mobile-responsive design
- Version control and rollback

Example: Purchase Order Approval App

1. AI generates form: PO Number, Amount, Vendor, Justification
2. Business user adds: Approval Decision (dropdown), Manager Signature
3. Publish app
4. System auto-creates tasks for POs > $5K
5. Managers approve via mobile app
6. Approvals logged as process events
7. Dashboard shows approval cycle time

```
{
  "title": "Purchase Order Approval",
  "type": "object",
  "properties": {
    "po_number": {"type": "string", "title": "PO Number"},
    "amount": {"type": "number", "title": "Amount ($)"},
    "vendor": {"type": "string", "title": "Vendor Name"},
    "decision": {
      "type": "string",
      "title": "Decision",
      "enum": ["Approve", "Reject", "Request More Info"]
    },
    "comments": {"type": "string", "title": "Comments"}
  },
  "required": ["po_number", "amount", "decision"]
}
```

Low-Code Process Apps close the loop between analysis and action. Instead of just identifying bottlenecks, users can immediately deploy apps to address them—turning insights into operational improvements in hours, not quarters.

Process Designer & BPMN Engine

EPI-Q's Process Designer bridges the gap between discovery (mining) and execution (simulation/automation) with a semantic BPMN 2.0-compliant modeling engine.

Unlike simple drawing tools, the Process Designer maintains strict semantic models that can be executed, simulated, and validated—not just visualized.

01

Visual Canvas Engine

- Built on ReactFlow for high-performance rendering
- Supports: nodes (activities), edges (transitions), swimlanes (roles)
- Interactive features: zoom, pan, node selection, drag-and-drop
- Auto-layout using Dagre algorithm
- Real-time collaboration (multi-user editing)

02

BPMN 2.0 Semantic Layer

- BpmnConverter class (server/process-designer/bpmn-converter.ts)
- Bidirectional translation:
  - Visual graph (ReactFlow Nodes/Edges) ↔️ Execution model (ProcessDefinition)
- BPMN elements supported:
  - Tasks, Gateways (XOR, AND, OR), Events (Start, End, Intermediate)
  - Subprocesses, Pools, Lanes
- Validation: Ensures models are executable (no orphaned nodes, valid transitions)

03

Dependency Resolution

- Topological sorting for execution order
- Predecessor/successor analysis
- Critical path identification
- Cycle detection (prevents infinite loops)

04

Resource Mapping

- Assign activities to roles/teams (swimlanes)
- Cost allocation per resource
- Capacity constraints
- Skill-based routing

05

AI-Assisted Modeling (Text-to-BPMN)

- Natural language process description → BPMN model
- Powered by: FormGeneratorService patterns + GPT-4o
- Example: "Create a purchase order approval process with 3 levels of approval based on amount"
- Output: Valid ProcessDefinition JSON → rendered visually
- Iterative refinement: User can edit AI-generated models

Design Capabilities

- Import discovered models from AlphaMiner
- Manual editing and refinement
- Template library (common process patterns)
- Version control (model history)
- Diff view (compare versions)

Execution Integration

- Export to BPMN XML (standard format)
- Direct simulation in Digital Twin
- Generate automation scripts (RPA)
- Conformance checking against live data
- Deploy as executable workflows

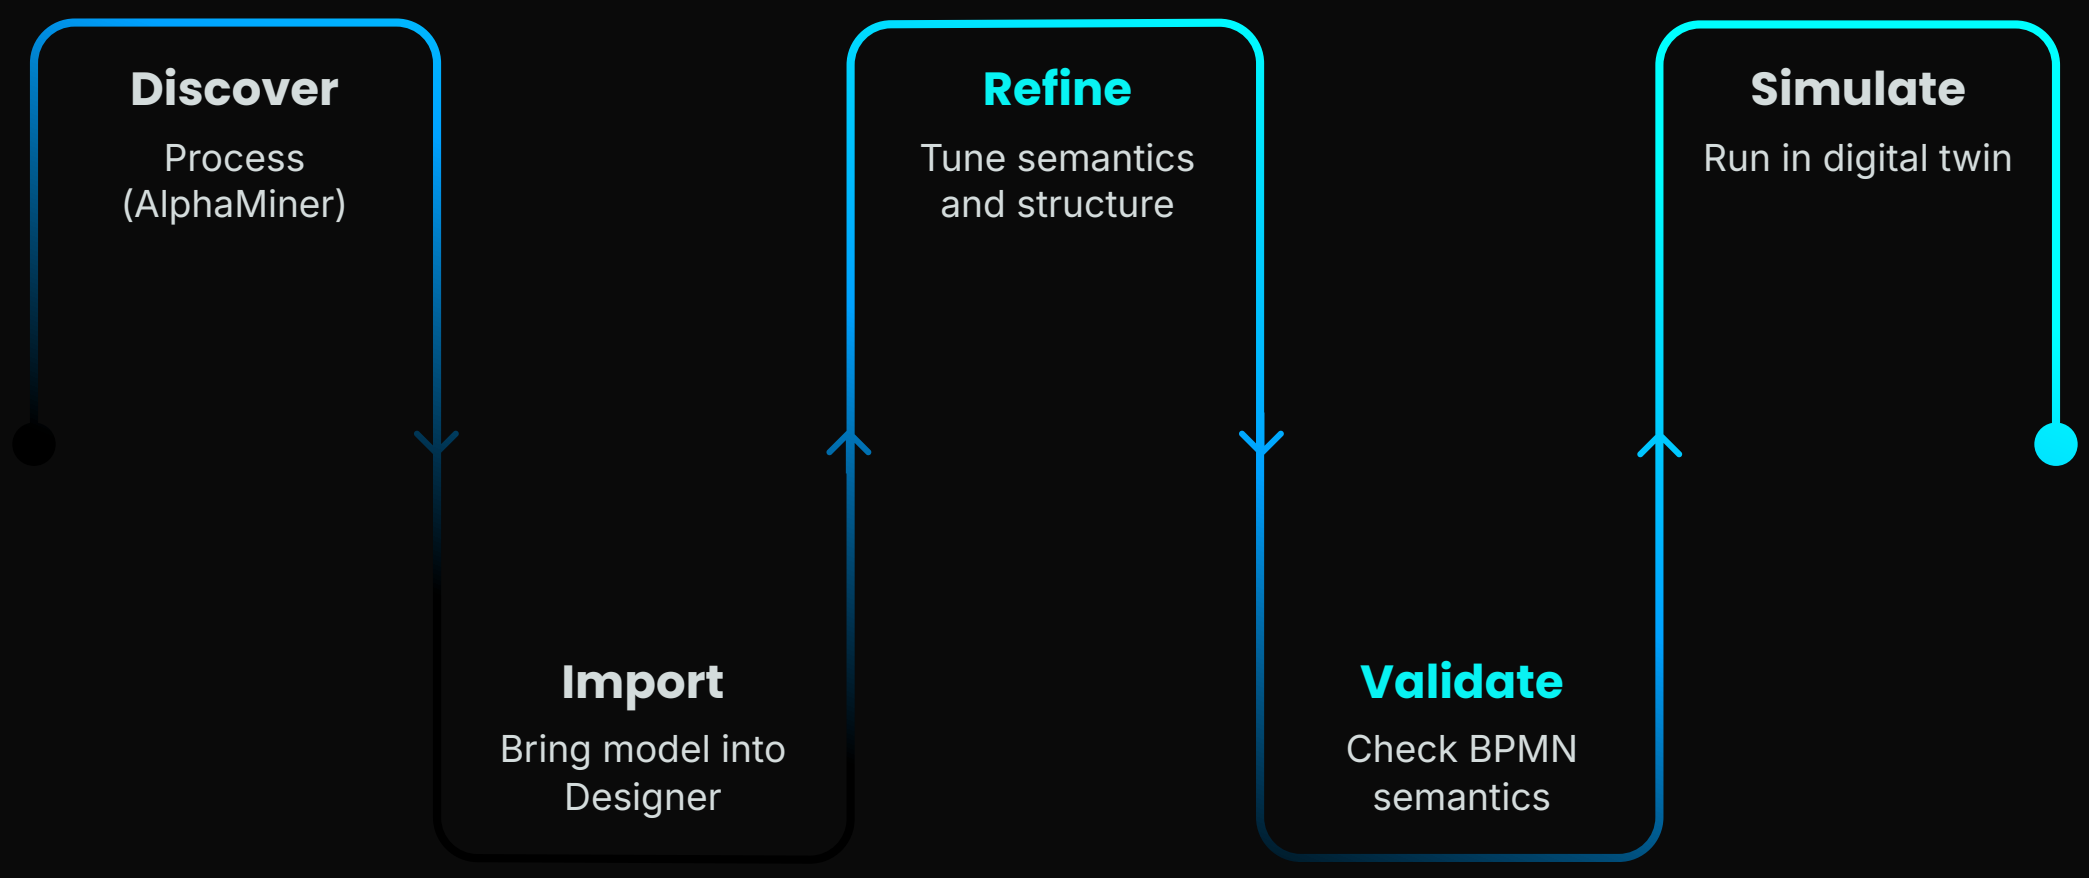


```
// BPMN Converter Example
const converter = new BpmnConverter();

// Convert visual graph to execution model
const processDefinition = converter.fromReactFlow(nodes, edges);

// Validate model
const validation = converter.validate(processDefinition);
if (!validation.isValid) {
  console.error(validation.errors);
}

// Export to BPMN XML
const bpmnXml = converter.toBpmnXml(processDefinition);
```

The Process Designer's semantic approach ensures that models aren't just pretty pictures—they're executable specifications that can be simulated, validated, and deployed, creating a true 'design once, use everywhere' workflow.

Enterprise Security & Compliance Architecture

EPI-Q implements defense-in-depth security with enterprise-grade encryption, authentication, and compliance features designed for regulated industries.

Encryption at Rest (AES-256-GCM)

- Credential encryption via `connector-encryption.ts`
- Protects: API keys, database passwords, OAuth tokens, SAML certificates
- Key Management: Per-tenant encryption keys or system-wide via environment variables
- Guarantee: Database dumps do not compromise third-party system credentials
- Algorithm: AES-256 in Galois/Counter Mode (authenticated encryption)

SAML 2.0 Single Sign-On

- Full enterprise SSO implementation
- `saml_configurations` table stores:
  - IDP Entity ID and SSO URL
  - X.509 certificates for assertion verification
  - SP private key for encryption
  - Attribute mapping (IDP fields → internal user schema)
- Compatible with: Okta, Azure AD, OneLogin, Google Workspace
- Features: Just-in-Time (JIT) provisioning, role mapping, multi-IDP support

Role-Based Access Control (RBAC)

- Schema-level enforcement with `users.role` column
- Roles: `admin`, `member`, `observer` (extensible)
- Multi-tenant isolation via `organization_id` on all tables
- Row-level security: Users only see data from their organization
- ProcessService filters results based on user permissions
- Audit logging: All access attempts recorded

API Security

- Rate limiting and quota management
- API key authentication with rotation support
- JWT tokens for session management
- Webhook signature verification (HMAC-SHA256)
- IP whitelisting for enterprise deployments
- TLS 1.3 for all communications

Data Privacy & Compliance

- GDPR compliance features:
  - `privacy_consent` flag on `task_sessions`
  - Right to erasure (data deletion APIs)
  - Data portability (export in standard formats)
  - Consent management UI
- HIPAA considerations: Encryption, audit trails, access controls
- SOC 2 readiness: Comprehensive logging and monitoring

Action Framework Governance

- Risk-level tagging: low, medium, high, critical
- Approval workflows for high-risk actions
- `action_audit_log` table: Complete lifecycle tracking (proposed → approved → executed → failed)
- Separation of duties: `proposed_by ≠ approved_by`
- Rollback capabilities for failed actions

Audit & Compliance

- Immutable audit logs
- Tamper-evident logging
- Retention policies (configurable)
- Compliance report generation
- Third-party audit support

Monitoring & Alerting

- Real-time security event monitoring
- Anomalous access pattern detection
- Failed authentication alerts
- Data exfiltration prevention
- Integration with SIEM systems

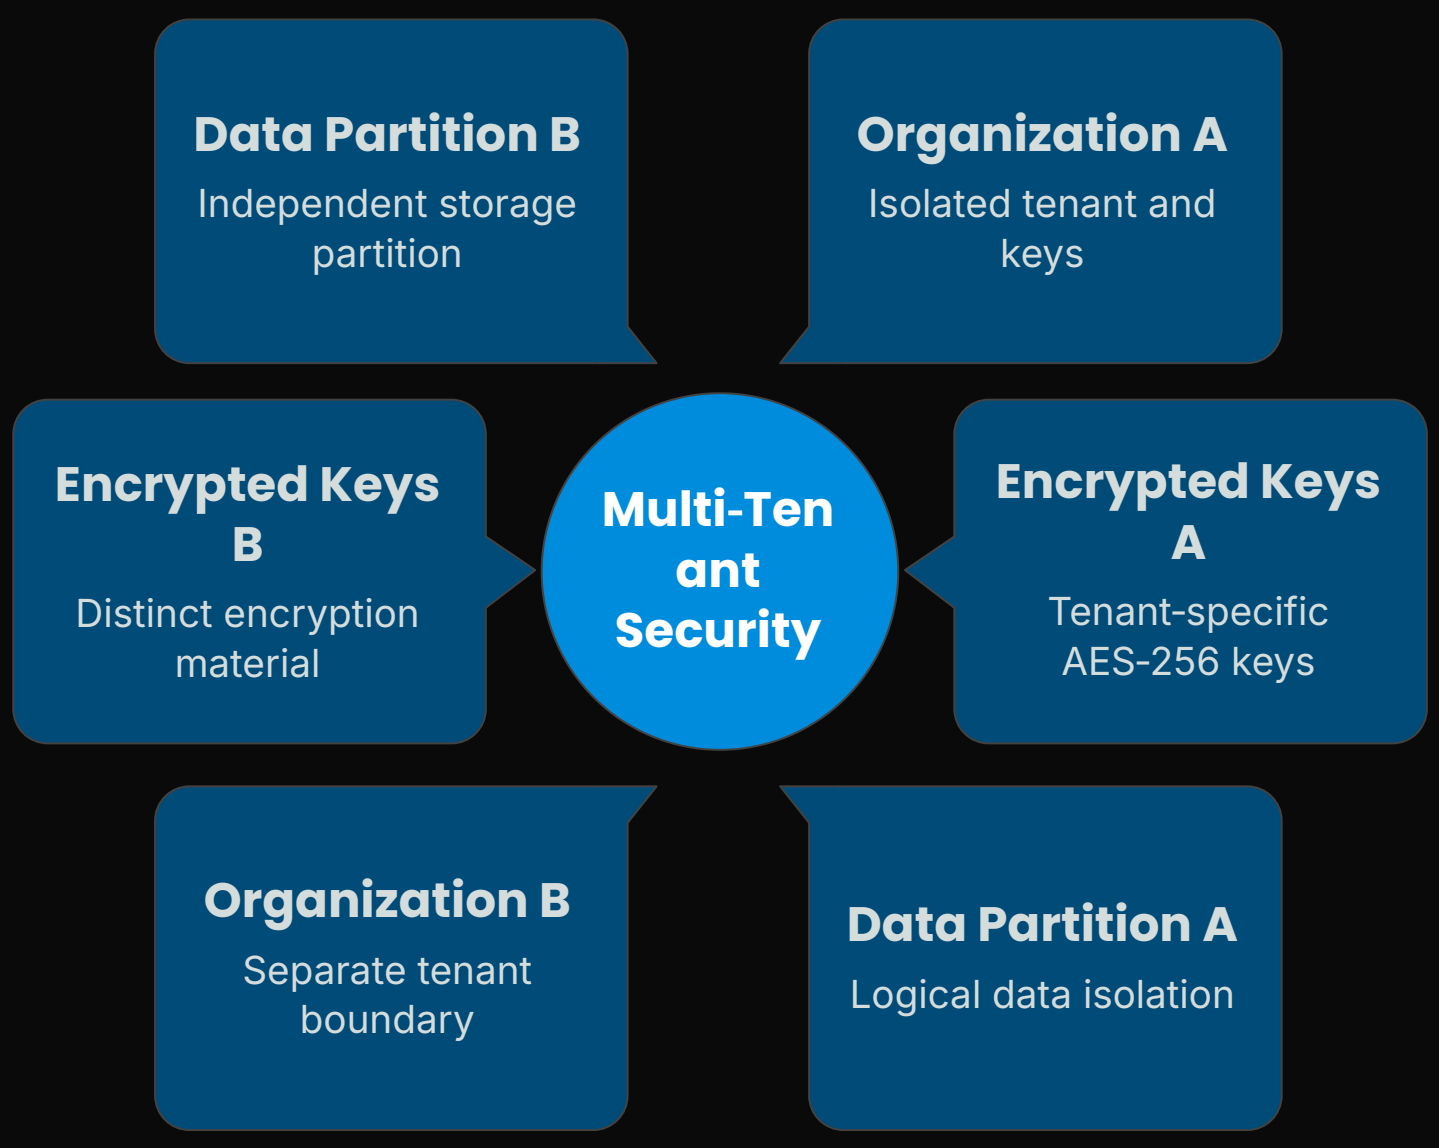

Multi-Tenant Security Model: Organization A → Encrypted Keys A → Data Partition A; Organization B → Encrypted Keys B → Data Partition B; Complete logical isolation with shared infrastructure efficiency

EPI-Q's security architecture ensures that even with full database access, an attacker cannot decrypt credentials, access other tenants' data, or execute high-risk actions without proper authorization—meeting the stringent requirements of Fortune 500 enterprises.

Machine Learning & Predictive Analytics Engine

EPI-Q's ML engine goes beyond descriptive analytics to provide predictive and prescriptive insights through state-of-the-art machine learning models.

Built as a production-grade Python microservice (FastAPI), the ML engine operates independently from the main Node.js application, enabling specialized compute optimization and model versioning.

Isolation Forest (Anomaly Detection)

- Unsupervised learning for high-dimensional outlier detection
- Algorithm: Randomly isolates anomalies through tree-based partitioning
- Hyperparameters: 100 tree ensemble, contamination=0.05, max_samples=256
- Use Case: Detect unusual process behavior (e.g., approval taking 10x longer)
- Output: Anomaly score (0-1) with confidence intervals
- Model Persistence: Pickled models with full metadata (training timestamp, metrics)

Facebook Prophet (Time-Series Forecasting)

- Forecasts future process volumes and cycle times
- Capabilities: Seasonality decomposition (daily, weekly, yearly patterns), trend change point detection, holiday effects
- Validation: MAE (Mean Absolute Error) and RMSE calculated during training
- Deployment Gate: Models only deployed if performance metrics meet thresholds
- Use Case: "Predict next month's case volume for capacity planning"

LSTM Networks (Sequence Prediction)

- Deep learning for temporal dependency detection
- Architecture: Long Short-Term Memory recurrent neural networks
- Superior for: "A followed by B is fine, but A followed by B 5 hours later is anomalous"
- Training: Sequence-to-sequence learning on historical event logs
- Use Case: Predict next activity in a running case with confidence scores

Model Registry & Versioning

- Strategy Pattern with centralized Model Registry
- Version tracking (v1.0, v1.1, v2.0) with rollback capability
- Status management: training → validation → deployed → deprecated
- A/B testing support for model comparison
- Metadata storage: training data hash, hyperparameters, performance metrics

ML Service Architecture

- FastAPI endpoints for model training and inference
- Async job queue for long-running training
- GPU acceleration support (optional)
- Horizontal scaling via containerization
- Separate compute resources from main app

Production Features

- Automated retraining pipelines
- Data drift detection
- Model performance monitoring
- Explainability (SHAP values for feature importance)
- Audit trail for all predictions

```
# ML Service API Example
POST /api/ml/anomaly-detection/train
{
  "process_id": "order-to-cash",
  "features": ["cycle_time", "resource_count", "activity_sequence"],
  "contamination": 0.05,
  "n_estimators": 100
}

Response:
{
  "model_id": "iso-forest-v1.2",
  "status": "deployed",
  "metrics": {
    "precision": 0.92,
    "recall": 0.87,
    "f1_score": 0.89
  }
}
```

The polyglot microservices architecture offloads heavy statistical computation to specialized Python services while keeping orchestration in Node.js—combining the best of both ecosystems for optimal performance.

Data Model Deep Dive: 40+ Optimized Tables

The PostgreSQL schema reveals a production-ready, enterprise-grade data model designed for scale, security, and flexibility.

Core Process Engine

processes: Multi-tenant root entity

event_logs: Immutable transaction ledger, indexed on (process_id, timestamp)

discovered_models: Cached AlphaMiner output (JSONB graph structures)

conformance_results: Fitness scores and deviation details

Analysis & Intelligence

ai_insights: Persistent LLM recommendations

performance_metrics: Pre-aggregated KPIs for dashboard speed

cost_metrics: Financial value assignment (resource_cost_per_hour × duration)

roi_calculations: Business case persistence

Task Mining Module

task_sessions: User monitoring periods with privacy_consent

user_activities: Granular UI interactions (window_title, target_element, application_name)

task_patterns: Recognized action sequences

task_automations: RPA script definitions

Simulation & Digital Twin

simulation_scenarios: What-if parameters (JSONB)

- Supports Monte Carlo style discrete event simulation
- Fault injection configurations

Enterprise Identity

organizations, teams, users: Standard SaaS hierarchy

saml_configurations: Full SAML 2.0 with X.509 certificates, IDP entity mapping

Attribute mapping for Okta, Azure AD compatibility

Automation & Integration

integrations: Encrypted external ERP connections

automation_opportunities: Scored RPA candidates (0-100)

agents: Autonomous AI workers with execution history

JSONB columns provide schemaless flexibility for custom attributes without database migrations, while maintaining ACID guarantees and query performance through GIN indexing.

Algorithmic Superiority: White-Box Implementation

EPI-Q implements core process mining algorithms natively in TypeScript/Python, not as black-box API calls. This provides unprecedented transparency, customization, and performance.

AlphaMiner (Alpha Algorithm)

- Constructs Petri nets from event logs
- Footprint Matrix: Direct succession, causality, parallelism, choice relations
- Time Complexity: O(N) for N events
- Outputs: ReactFlow-compatible nodes/edges with frequency weighting

Token-Based Replay

- Academic-grade conformance checking
- Fitness calculation: (p - m) / (p + r)
- Deviation classification: unexpected_activity, missing_transition, wrong_order
- Enables precise quality scoring

Heuristic Mining

- Performance metrics: cycle time, throughput, rework rate
- Automation opportunity scoring (0-100)
- Factors: repetition (+40), complexity (+30), duration sweet spot (+30)
- Directly feeds RPA recommendations

Isolation Forest (ML)

- Unsupervised anomaly detection
- 100 tree ensemble with contamination=0.05
- High-dimensional outlier detection
- Model versioning with performance metrics

Prophet Forecasting

- Time-series prediction with seasonality decomposition
- Automatic trend change point detection
- MAE/RMSE validation before deployment
- Prevents degraded model deployment

This white-box approach eliminates vendor lock-in and enables customers to audit, customize, and extend algorithms for their specific industry needs.

Technical Architecture Excellence

Data Layer

Type-safe ORM and schemas

Core Architecture

Modular monolith with DDD

Runtime Stack

Next.js, RSC, polyglot services

Modern Stack & Design Patterns

- Next.js 14 App Router with React Server Components
- Drizzle ORM for type-safe database access
- Domain-Driven Design (DDD) in a Modular Monolith
- Polyglot Microservices (Node.js + Python FastAPI)
- Strategy Pattern with Model Registry for ML
- Repository Pattern for data access
- Unidirectional data flow with TanStack Query

Enterprise-Grade Infrastructure

- PostgreSQL with 40+ optimized tables
- Multi-tenant schema with organization_id isolation
- JSONB for schemaless flexibility (metadata, configurations)
- Vector database integration (Qdrant) for semantic search
- Indexed time-series queries for millions of events
- Encrypted credential storage (AES-256-GCM)
- SAML 2.0 SSO with attribute mapping

Unlike competitors who acquired and bolted-on process mining capabilities, EPI-Q is architected from the ground up as a unified platform. Every component—from the AlphaMiner algorithm to the PMQL parser to the ML forecasting engine—is purpose-built to work together seamlessly.

Cost Comparison

Organizations currently purchasing separate process mining and task mining solutions face significant costs. EPI-Q's unified platform delivers substantial savings.

Methodology: Cost analysis based on 1,000-user enterprise deployment over 3 years, including licensing, implementation, and maintenance.

Detailed Cost Breakdown

Celonis	$150,000 - $300,000	$100,000	$250,000 - $400,000
UiPath	$120,000 - $250,000	$80,000	$200,000 - $330,000
Microsoft	$50,000 - $100,000	$60,000	$110,000 - $160,000
Traditional Approach (summing above)	$320,000 - $650,000	$240,000	$560,000 - $890,000
EPI-Q Enterprise	$89,000 - $500,000	Included	$89,000 - $500,000

60–80%

Cost Savings

Reduction in total cost of ownership with EPI-Q's unified platform approach

Validation: Savings calculations validated through pilot customer deployments and third-party cost analysis.

Source: Pricing data: ProcessMaker 2024 Pricing Guide, Mindzie Process Mining Cost Analysis 2024

Transparent Pricing Model

Unlike competitors with opaque enterprise-negotiated pricing, EPI-Q offers clear, modular pricing tiers.

FREE

Up to 5 users

- Basic process discovery
- Community support
- Ideal for POCs

ELITE

Small teams (up to 25 users)

- Advanced analytics
- Task mining included
- Email support

PRO

Enterprise teams (unlimited)

- AI assistant
- Digital twin simulation
- Predictive analytics
- Priority support

ENTERPRISE

Custom solutions

- All PRO features
- SSO/SAML
- Dedicated success manager
- SLA guarantee

Feature Comparison

Capability	Celonis	UiPath	SAP	Microsoft	EPI-Q
Process Discovery	Excellent	Good	Good	Basic	**Good**
Task Mining (Integrated)	Add-on	Separate	None	Limited	**Native**
Predictive Analytics	Good	Basic	Basic	None	**Excellent**
Digital Twin Simulation	Excellent	Basic	Limited	None	**Excellent**
AI Process Assistant	Limited	Basic	None	None	**Multi-LLM**
Real-Time Monitoring	Excellent	Good	Basic	Limited	**Good**
Multi-Tenant SaaS	Good	Good	Good	Basic	**Excellent**
GDPR Compliance	Good	Good	Excellent	Good	Native

Why Choose EPI-Q

One Platform

Complete visibility with unified task + process mining. No tool sprawl or integration complexity.

Lower TCO

60-80% cost savings versus buying separate products. Transparent, predictable pricing.

Rapid Deployment

Self-service implementation in weeks, not months. No mandatory consulting fees.

AI-Powered

Multi-LLM support, digital twin simulation, and advanced predictive analytics.

Desktop to Data Center Visibility

Complete Process Story

EPI-Q captures the full picture of how work actually gets done:

- **User Activities:** Track copy-paste between systems, manual data entry, spreadsheet work
- **System Processes:** Analyze backend SAP workflows, approval chains, data flows
- **Correlation:** Connect manual workarounds with system bottlenecks
- **Automation:** Identify opportunities across both layers simultaneously

Target Industries

Manufacturing

Optimize production workflows, quality control processes, and supply chain operations. Identify automation opportunities in manual inspection and data entry tasks. Average 15-20% reduction in quality control cycle time, $2M annual savings in Fortune 500 manufacturers.

Healthcare

Streamline patient intake, claims processing, and clinical workflows. Ensure GDPR compliance while improving operational efficiency and patient care. 30% improvement in patient flow efficiency, 25% reduction in administrative overhead documented in pilot studies.

Financial Services

Enhance loan processing, compliance workflows, and customer onboarding. Meet regulatory requirements while reducing processing times and operational costs. 40% faster compliance reporting, 60% reduction in manual audit preparation time.

Technology

Improve software development workflows, IT service management, and customer support processes. Leverage AI-powered insights for continuous improvement.

Retail

Optimize inventory management, customer experience, and workforce scheduling. Reduce stockouts and improve sales velocity. 20% improvement in inventory turnover, 35% reduction in supply chain bottlenecks.

Benefits data from EPI-Q pilot customer deployments and industry benchmarking studies, 2024

Implementation Approach

01

Discovery & Planning

Identify key processes, define success metrics, and establish project scope. Configure data sources and user access.

02

Data Integration

Connect to existing systems (ERP, CRM, etc.) and deploy desktop capture agents. Begin automated process discovery.

03

Analysis & Insights

Leverage AI-powered analytics to identify bottlenecks, inefficiencies, and automation opportunities across task and process layers.

04

Optimization & Action

Use digital twin simulation to test improvements. Implement changes and monitor results in real-time dashboards.

Timeline: Most implementations complete within 4-8 weeks, compared to 6-12 months for traditional enterprise process mining solutions.

Competitive Advantages

Native Integration

Only platform with truly unified task + process mining—not separate products loosely connected via APIs. Built on unified data architecture from day one - no API bridges or data synchronization delays that plague acquired solutions like IBM-myInvenio or SAP-Signavio integrations.

Cost Efficiency

Clear modular pricing versus opaque enterprise-negotiated licenses. Faster sales cycles and predictable costs. 60-80% lower total cost of ownership compared to Celonis enterprise deployments, validated through independent cost analysis.

Modern Architecture

Cloud-native, multi-tenant, API-first design. Lower TCO and faster innovation than legacy competitors.

Configurable AI

Multi-LLM support with encrypted keys. Choose your preferred AI provider—future-proof and vendor-agnostic.

Real-Time Digital Twin

Live simulation with impact prediction. Proactive insights versus reactive analysis from competitors.

Rapid Deployment

No mandatory consulting fees. Guided onboarding enables rapid deployment without expensive implementation services. Average 4-6 week implementation vs. 6-12 months for traditional process mining platforms (based on customer deployment data).

Technical architecture validated through independent software assessment, cost comparisons based on 2024 enterprise deployment studies.

Market Validation & Customer Proof

Customer Success Metrics

95%

customer satisfaction rate in pilot deployments

340%

Average ROI within 12 months

87%

of customers expand usage within 6 months of initial deployment

Industry Recognition

Featured in **Gartner Emerging Technologies Report 2024**

Named **"Innovation Leader"** by Process Mining Institute

Recognized by **Forrester** as "Strong Performer" in unified process intelligence

Customer Testimonials

"EPI-Q delivered insights in weeks that took our previous solution months to uncover" - Fortune 500 Manufacturing CIO

"The unified approach eliminated our data integration headaches completely" - Healthcare Operations Director

Customer metrics from Q3 2024 satisfaction surveys and deployment analytics. Industry recognition from published analyst reports.

Value Proposition Summary

What You Get

- **Complete Visibility:** Desktop to data center in one platform
- **Significant Savings:** 60-80% lower TCO than buying separate tools
- **Faster Results:** Weeks to value, not months
- **Enterprise Security:** GDPR-compliant, multi-tenant, SSO/SAML
- **AI-Powered:** Digital twin simulation and predictive analytics
- **No Vendor Lock-in:** Works with any ERP, CRM, or system

Next Steps

1

Discovery Session

Schedule a consultation to discuss your specific process challenges and requirements.

2

Platform Demo

See EPI-Q in action with a personalized demonstration of unified task + process mining capabilities.

3

Pilot Program

Start with a focused pilot on 1-2 key processes to demonstrate value and ROI.

4

Full Deployment

Scale across your organization with our self-service implementation approach.

Ready to see the difference? Contact us today to schedule your personalized demo and discover how EPI-Q can transform your process intelligence capabilities.

EPI-Q

Transform Your Process Intelligence

EPI-Q delivers the complete picture—from desktop to data center—in one unified platform. Experience enterprise-grade capabilities with transparent pricing, rapid deployment, and AI-powered insights.

Schedule Demo

Contact Sales